

PLA architecture — frozen vision-language backbone fused with learned proximity tokens, decoded by an Action Chunking Transformer.

